

Week 3 Tutorial Sheet

(To be completed during the Week 3 tutorial class)

Objectives: The tutorials, in general, give practice in problem solving, in analysis of algorithms and data structures, and in mathematics and logic useful in the above.

Instructions to the class: Aim to attempt these questions before the tutorial! It will probably not be possible to cover all questions unless the class has prepared them in advance. There are marks allocated towards active participation during the class. You **must** attempt the problems under **Assessed Preparation** section **before** your tutorial class and give your worked out solutions to your tutor at the start of the class – this is a hurdle and failing to attempt these problems before your tutorial will result in 0 mark for that class even if you actively participate in the class.

Instructions to Tutors:

1. The purpose of the tutorials is not to solve the practical exercises!
2. The purpose is to check answers, and to discuss particular sticking points, not to simply make answers available.

Supplementary problems: The supplementary problems provide additional practice for you to complete after your tutorial class, or as pre-exam revision. Problems that are marked as **(Advanced)** difficulty are beyond the difficulty that you would be expected to complete in the exam, but are nonetheless useful practice problems as they will teach you skills and concepts that you can apply to other problems.

Assessed Preparation

Problem 1. Write pseudocode for insertion sort, except instead of sorting the elements into non-decreasing order, sort them into non-increasing order. Identify a useful invariant of this algorithm.

Problem 2. Consider the following algorithm that returns the number of occurrences of *target* in the sequence *A*. Identify a useful invariant that is true at the beginning of each iteration of the **while** loop. Prove that it holds, and use it to prove that the algorithm is correct.

```
1: function COUNT(A[1..n], target)
2:   count = 0
3:   i = 1
4:   while i ≤ n do
5:     if A[i] = target then
6:       count = count + 1
7:     end if
8:     i = i + 1
9:   end while
10:  return count
11: end function
```

Tutorial Problems

Problem 3. What are the complexities of these two approaches to stable counting sort outlined in lecture 2? What are the advantages and disadvantages of each?

Problem 4. Consider the following algorithm that returns the minimum element of a given sequence *A*. Identify

a useful invariant that is true at the beginning of each iteration of the **for** loop. Prove that it holds, and use it to show that the algorithm is correct.

```

1: function MINIMUM_ELEMENT( $A[1..n]$ )
2:    $\text{min} = A[1]$ 
3:   for  $i = 2$  to  $n$  do
4:     if  $A[i] < \text{min}$  then
5:        $\text{min} = A[i]$ 
6:     end if
7:   end for
8:   return  $\text{min}$ 
9: end function

```

Problem 5. Describe a simple modification that can be made to any comparison-based sorting algorithm to make it stable. How much space and time overhead does this modification incur?

Problem 6. Write an in-place algorithm that takes a sequence of n integers and removes all duplicate elements from it. The relative order of the remaining elements is not important. Your algorithm should run in $O(n \log(n))$ time and use $O(1)$ auxiliary space (i.e. it must be in-place).

Problem 7. Think about and discuss with those around you why auxiliary space complexity is a useful metric. Why is it often more informative than total space complexity? For the purpose of this problem, define auxiliary space complexity as the amount of space required by an algorithm, excluding the space taken by the input, and define total space complexity as the space taken by an algorithm, including the space taken by the input.

Problem 8. Devise an efficient online algorithm¹ that finds the smallest k elements of a sequence of integers. Write pseudocode for your algorithm. [Hint: Use a data structure that you have learned about in a previous unit]

Problem 9. Write an iterative Python function that implements binary search on a sorted, non-empty list, and returns the position of the key, or None if it does not exist.

- If there are multiple occurrences of the key, return the position of the **final** one. Identify a useful invariant of your program and explain why your algorithm is correct
- If there are multiple occurrences of the key, return the position of the **first** one. Identify a useful invariant of your program and explain why your algorithm is correct

Problem 10. A subroutine used by Mergesort is the merge routine, which takes two sorted lists and produces from them a single sorted list consisting of the elements from both original lists. In this problem, we want to design and analyse some algorithms for merging many lists, specifically $k \geq 2$ lists.

- Design an algorithm for merging k sorted lists of total size n that runs in $O(nk)$ time
- Design a better algorithm for merging k sorted lists of total size n that runs in $O(n \log(k))$
- Is it possible to write a comparison-based algorithm that merges k sorted lists that is faster than $O(n \log(k))$?

Supplementary Problems

Problem 11. Consider the problem of finding a target value in sequence (not necessarily sorted). Given below is pseudocode for a simple linear search that solves this problem. Identify a useful loop invariant of this algorithm and use it to prove that the algorithm is correct.

```

1: function LINEAR_SEARCH( $A[1..n]$ , target)

```

¹In this case, online means that you are given the numbers one at a time, and at any point you need to know which are the smallest k

```

2:   Set index = null
3:   for i = 1 to n do
4:       if A[i] = target then
5:           index = i
6:       end if
7:   end for
8:   return index
9: end function

```

Problem 12. Devise an algorithm that given a sorted sequence of distinct integers a_1, a_2, \dots, a_n determines whether there exists an element such that $a_i = i$. Your algorithm should run in $O(\log(n))$ time.

Problem 13. Consider the following variation on the usual implementation of insertion sort.

```

1: function FAST_INSERTION_SORT(A[1..n])
2:   for i = 2 to n do
3:       Set key = A[i]
4:       Binary search to find max  $k < i$  such that  $A[k] \leq \text{key}$ 
5:       for j = k + 1 to i do
6:            $A[j] = A[j - 1]$ 
7:       end for
8:        $A[k] = \text{key}$ 
9:   end for
10: end function

```

- What is the number of comparisons performed by this implementation of insertion sort?
- What is the worst-case time complexity of this implementation of insertion sort?
- What do the above two facts imply about the use of the comparison model (analysing a sorting algorithm's complexity by the number of comparisons it does) for analysing time complexity?

Problem 14. (Advanced) Consider the problem of integer multiplication. For two integers with n digits, the standard multiplication algorithm that you learned in school runs in $O(n^2)$ time. Interestingly, we can do much better than this. We will consider integers written in binary and suppose that n is a power of 2, but the techniques described here work in any base. Let's analyse a divide-and-conquer approach in which we split each of the integers into two halves and recursively multiply the two halves. Given two integers X and Y of n bits each, we write X_L, X_R, Y_L, Y_R for the left and right halves of X and Y respectively. X and Y are decomposed as follows

$$X = (2^{\frac{n}{2}} X_L + X_R), \quad Y = (2^{\frac{n}{2}} Y_L + Y_R),$$

and hence we can write their product as

$$XY = (2^{\frac{n}{2}} X_L + X_R)(2^{\frac{n}{2}} Y_L + Y_R) = 2^n X_L Y_L + 2^{\frac{n}{2}} (X_L Y_R + X_R Y_L) + X_R Y_R. \quad (1)$$

Suppose we implement a divide-and-conquer algorithm that uses this formula directly to compute the product. Note that multiplication by a power of two takes $O(n)$ time since this is just a shift operation, and keep in mind that adding two n bit numbers takes $O(n)$ time.

- Write a recurrence relation for the time complexity of such an algorithm
- Solve the recurrence relation and write the time complexity in big-O notation

The above method requires four recursive calls since there are four products to compute in the given formula. Suppose instead that we make use of the following formula

$$(X_L Y_R + X_R Y_L) = (X_L + X_R)(Y_R + Y_L) - X_R Y_R - X_L Y_L.$$

If we use this formula for evaluating Equation (1), we can reduce the required number of multiplications to three, so an algorithm based on this formula should be faster.

- (c) Write a recurrence relation for the time complexity of such an algorithm
- (d) Solve the recurrence relation and write the time complexity in big-O notation
- (e) Implement this algorithm in Python. Compare it against the naive $O(n^2)$ algorithm for very large numbers and measure the difference in running time

This algorithm is known as Karatsuba multiplication and is what Python uses to multiply very large numbers. Faster algorithms based on the Fast Fourier Transform exist (see the Schönhage–Strassen algorithm), but these are rarely used in practice due to having very large constant factors.

Problem 15. (Advanced) When we analyse the complexity of an algorithm, we always make assumptions about the kinds of operations we can perform, how long they will take, and how much space that we will use. We call this the *model of computation* under which we analyse the algorithm. The assumptions that we make can lead to wildly different conclusions in our analysis.

An early model of computation used by computer scientists was the *RAM*, the Random-Access Machine. In the RAM model, we assume that we have unlimited memory, consisting of *registers*. Each register has an *address* and some contents, which can be any integer. Additionally, integers are used as pointers to refer to memory addresses. A RAM is endowed with certain operations that it is allowed to perform in constant time. The total amount of space used by an algorithm is the total size of all of the contents of the registers used by it.

- (a) State some unrealistic aspects of the RAM model of computation
- (b) Explain why the definition of an in-place algorithm being those which use $O(1)$ auxiliary space is near worthless in this description of the RAM model
- (c) In the Week 2 tutorial, we discussed fast algorithms for computing Fibonacci numbers. In particular, we saw that $F(n)$ can be computed using matrix powers, which can be computed in just $O(\log(n))$ multiplications. If a RAM is endowed with all arithmetic operations, then we can compute $F(n)$ in $O(\log(n))$ time using this algorithm. If instead the RAM is only allowed to perform addition, subtraction, and bitwise operations in constant time, we can still simulate multiplication using any multiplication algorithm (for example, Karatsuba multiplication shown in Problem 14). Explain why in this model, it is impossible to compute $F(n)$ in $O(\log(n))$ time with this algorithm.

A model that is more commonly used in modern algorithm and data structure analysis is the *word RAM* (short for word Random-Access Machine). In the word RAM model, every word is a fixed-size w -bit integer, where w is a parameter of the model. We can perform all standard arithmetic and bitwise operations on w -bit integers in constant time. The total amount of space used by an algorithm is the number of words that it uses

- (d) What is the maximum amount of memory that can be used by a w -bit word RAM?
- (e) Suppose we wish to solve a problem whose input is a sequence of size n . What assumption must be made about the model for this to make sense?
- (f) Discuss some aspects that the word RAM still fails to account for in realistic modern computers
- (g) Does the word RAM model allow us to compute $F(n)$, the n^{th} Fibonacci number, in $O(\log(n))$ time?