

Faculty of Information Technology, Monash University

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

This material has been reproduced and communicated to you by or on behalf of Monash University pursuant to Part VB of the Copyright Act 1968 (the Act). The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act. Do not remove this notice

FIT2004: Algorithms and Data Structures

Week 12: Topological Sort and Design Principles

These slides are prepared by [M. A. Cheema](#) and are based on the material developed by [Arun Konagurthu](#) and [Lloyd Allison](#).

Announcements/Things to note

- Complete SETU before it closes (21 June)
- The post-semester consultation schedule is the same as the in semester consultation schedule
- Any changes to the schedule will be announced on Moodle

Overview


- Topological Sort
 - Kahn's Algorithm
 - Depth First Search
- Design Principles (FIT2004 Summary)
- Final Exam etc.
- Review of all lecture material

Directed Acyclic Graph (DAG)


A Directed Acyclic Graph (DAG) is

- Directed
- Acyclic – has no cycles
- Graph

Which of the two graphs is a DAG?





Graph 1



Graph 2

DAG: Examples

- sub-tasks of a project and which “must finish before”
 - $A \rightarrow B$ means task A must finish before task B
 - so, DAGs useful in project management
- relationships between subjects for your degree -- “is prerequisite for”
 - $A \rightarrow B$ means subject A must be completed before enrolling in subject B
- people genealogy – “is an ancestor of”
 - $A \rightarrow B$ means A is an ancestor of B
- power sets and “is a subset of”
 - $A \rightarrow B$ means A is a subset of B



[Source: wikipedia](#)


Topological Sort of a DAG

Order of vertices in a DAG

- $A < B$ if $A \rightarrow B$.
 - Note that if $A \rightarrow B$ and $B \rightarrow D$, we have $A < B$ and $B < D$ which implies that $A < D$ (i.e., transitivity).
- Some vertices may be incomparable (e.g., B and C are incomparable), i.e. $A < B$ and $A < C$ but we do not know whether $C < B$ or $B < C$.

A topological order

- is a permutation of the vertices in the original DAG such that
- for **every** directed edge $u \rightarrow v$ of the DAG
 - ✦ u appears before v in the permutation



Example: A, B, C, E, D

- Topological sort of a DAG of “is prerequisite of” example gives an ordering of the subjects for studying your degree, one at a time, while obeying prerequisite rules.

Topological Sort of a DAG



- A DAG can have many valid topological sorts, e.g., let u and v be two incomparable vertices, u may appear before or after v .

Which of these is NOT a valid topological ordering of the DAG

1. A, B, C, E, D
2. A, C, B, E, D
3. A, C, E, B, D
4. A, B, E, C, D

How to do topological sort?

- Kahn's Algorithm




Overview

- Topological Sort
 - Kahn's Algorithm
 - Depth First Search
- Design Principles (FIT2004 Summary)
- Final Exam etc.
- Review of all lecture material

Kahn's Algorithm: High level idea


For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v



Sorted:

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v




Sorted:

A

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v




Sorted:


A

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v




Sorted:



Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v




Sorted:

A	B
----------	----------

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v




Sorted:

A	B	C
---	---	---

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v




Sorted:

A	B	C
---	---	---

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v




Sorted:

A	B	C	E
---	---	---	---

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v



Sorted:


A	B	C	E
---	---	---	---

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge

Add v to sorted

Remove the outgoing edges of v



Sorted:

A	B	C	E	D
---	---	---	---	---

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge


Add v to sorted

Remove the outgoing edges of v

How can we efficiently track the number of incoming edges?

Quiz time!

<https://flux.qa> - RFIBMB



Sorted:

A	B	C	E	D
---	---	---	---	---


Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge

Add v to sorted

Remove the outgoing edges of v

How can we efficiently track the number of incoming edges?



A	B	C	D	E
0	1	1	3	1

Order:

Sorted:

A	B	C	E	D
---	---	---	---	---

Kahn's Algorithm: High level idea


For each vertex v that does not have ANY incoming edge

Add v to sorted

Remove the outgoing edges of v

How can we efficiently track the number of incoming edges?

When we remove A , update it's children by -1



A	B	C	D	E
0	1	1	3	1

Order:

A	B	C	E	D
---	---	---	---	---

Sorted:

Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge

Add v to sorted

Remove the outgoing edges of v

How can we efficiently track the number of incoming edges?


When we remove A , update it's children by -1

Order:

A	B	C	D	E
0	0	1	3	1

Sorted:

A	B	C	E	D
---	---	---	---	---



Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge

Add v to sorted

Remove the outgoing edges of v

How can we efficiently track the number of incoming edges?


When we remove A , update it's children by -1

Order:

A	B	C	D	E
0	0	0	3	1

Sorted:

A	B	C	E	D
---	---	---	---	---



Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge

Add v to sorted

Remove the outgoing edges of v

How can we efficiently track the number of incoming edges?

When we remove A , update it's children by -1

Complexity of such an approach?

Quiz time!


<https://flux.qa> - RFIBMB

A	B	C	D	E
0	0	0	3	1

Order:

A	B	C	E	D
---	---	---	---	---

Sorted:



Kahn's Algorithm: High level idea

For each vertex v that does not have ANY incoming edge
Add v to sorted
Remove the outgoing edges of v

- Loop occurs V times
 - Finding a vertex with 0 in “order” takes $O(V)$
 - Adding to sorted is $O(1)$
 - Removing an outgoing edge costs $O(1)$ (using order)
 - This happens E times over the life of the algorithm
- So this algorithm would be $O(V*V + E) = O(V^2)$
- Can we do better?

Quiz time!

<https://flux.gq> - RFIBMB

Sorted:

A	B	C	E	D
---	---	---	---	---

Order:

A	B	C	D	E
0	0	0	3	1

Kahn's Algorithm: Detailed pseudocode

```
1: function TOPOLOGICAL_SORT( $G = (V, E)$ )
2:    $order$  = empty array
3:    $in\_degree$  = array of size  $V$ , initialised with the number of incoming edges to each vertex
4:    $ready$  = queue of all vertices with no incoming edges
5:   while  $ready$  is not empty do
6:      $u$  =  $ready.pop()$ 
7:      $order.append(u)$ 
8:     for each edge  $(u, v)$  adjacent to  $u$  do
9:        $in\_degree[v] -= 1$ 
10:      if  $in\_degree[v] = 0$  then
11:         $ready.push(v)$ 
12:   return  $order$ 
```

- Loop occurs V times
 - Pop is $O(1)$, append is $O(1)$
 - Inner loop runs E times in total
 - Removing an edge is $O(1)$
 - If is $O(1)$, push is $O(1)$
- So this algorithm would be **$O(V + E)$**

Overview

- **Topological Sort**
 - Kahn's Algorithm
 - **Depth First Search**
- Design Principles (FIT2004 Summary)
- Final Exam etc.
- Review of all lecture material

Depth First Search (DFS)

Below is the DFS algorithm we saw in week 8

- **function DFS(v):**
 - Mark v as Visited
 - For each adjacent edge (v,u)
 - ✦ If u is not visited
 - DFS(u)


Assume we call DFS(A), which of the following is NOT a possible order in which vertices are marked visited.

A, B, D, C, E

A, C, E, D, B

A, C, D, E, B


A, C, E, B, D



DFS for Topological Sort



Algorithm 72 Topological sorting using DFS

```
1: function TOPOLOGICAL_SORT( $G = (V, E)$ )
2:    $order = \text{empty array}$ 
3:    $visited[1..n] = \text{false}$ 
4:   for each vertex  $v = 1$  to  $n$  do
5:     if not  $visited[v]$  then
6:       DFS( $v$ )
7:   return  $\text{reverse}(order)$ 
8:
9: function DFS( $u$ )
10:   $visited[u] = \text{true}$ 
11:  for each vertex  $v$  adjacent to  $u$  do
12:    if not  $visited[v]$  then
13:      DFS( $v$ )
14:   $order.append(u)$ 
```



*// Add to order **after** visiting descendants*

Sorted:



Overview

- Topological Sort
 - Kahn's Algorithm
 - Depth First Search
- Design Principles (FIT2004 Summary)
- Final Exam etc.
- Review of all lecture material

Design Principles (Summing up FIT2004)

Here are some broad strategies to (try to) solve algorithmic problems:

- Look out for good invariants to exploit
- Attempt to balance your work as much as possible
- Do not repeat work (so, store and re-use!)
- Use appropriate data structures
- Try well-known problem solving strategies
- Sometimes greed is good!
- These are general guidelines. As always, there are many exceptions

Look out for good invariants to exploit

- Here are **some** algorithms we considered in the unit that do precisely this!
- Binary Search (Refer Week 2 lecture)
- Sorting (Refer Lectures from Weeks 2 and 3)
- Shortest Paths and Connectivity
 - Dijkstra's algorithm (Refer Week 8 Lectures)
 - Floyd-Warshall algorithm (Refer Week 9 lectures)
- Minimum Spanning Tree Algorithms (Refer Week 10 lectures)

Balance your work as much as possible

- For problems that allow division of labour (eg. Divide and Conquer)
- Try to divide work **equally** as much as possible
- Merge sort achieves this
 - $O(N \log N)$ -time always!
- Quick sort does not necessarily achieve this – depends on the choice of the pivot (Refer week 3)
 - Good pivots give $O(N \log N)$ -time
 - Bad pivots give $O(N^2)$ -time

Choose Data Structures with care

- Certain data representations are more efficient than others for a given problem
- Priority Queue in Dijkstra's algorithm (Refer Week 8)
- Union-Find data structure in Kruskal's algorithm (refer Week 9)
- Efficient Search and retrieval data structures of various kinds (Refer Weeks 5,6,7 lectures)

Don't repeat work

- Do not compute anything more than once (if there is room to store it for reuse)
- Underpins Dynamic Programming strategy
 - Edit Distance (Refer Week 4 Lecture)
 - Knapsack Problem (Refer Week 4 Lecture)

Try well known problem solving strategies

- Divide and Conquer (Refer Weeks 3, 4 lectures)
- Dynamic Programming (Refer Weeks 4, 8, 9 lectures)

Sometimes greed is good

- A **greedy strategy** is to make a “local” choice based on current information
- Sometimes gives optimal solution, e.g.
 - Dijkstra’s single source shortest paths algorithm (Refer Week 8 Lectures)
 - Minimum Spanning Tree Algorithms – Prim’s and Kruskal’s (Refer Week 10 lectures) minimum spanning tree algorithm.
- **Greedy is sometimes a good heuristic!**
 - Sometimes gives a “good” solution to a (combinatorial) problem even if not guaranteed optimal

Overview

- Topological Sort
 - Kahn's Algorithm
 - Depth First Search
- Design Principles (FIT2004 Summary)
- Final Exam etc.
- Review of all lecture material

Final Exam

- Time allowed: 2 hours + 10 minutes reading time
- Total Marks: 60
- Exam is open book
 - If a question asks you to **describe an algorithm**, you can write your idea in plain English
 - If a question asks you to **write pseudocode**, you **must** write your idea in a more structured way (like the ones in lecture slides or even Python code)
 - If a question asks for complexity, it means **big-O**, the **tightest bound** and the **worst case** unless otherwise specified
- Hurdles:
 - At least 16 out of 40 marks in in-semester assessments (assignment + mid-semester test + lecture/tutorial participation)
 - At least 24 out of 60 marks in the final exam
 - At least 50 marks overall
- Do not miss final exam even if you fail in-semester hurdle.
 - It affects your WAM

Non-Examinable Content

- Additional material in lecture notes is NOT examinable
 - In other words, anything NOT covered in lectures, tutorials, labs is NOT EXAMINABLE!
- Advanced questions in tutorials are NOT examinable
- Anything marked “not examinable” is not examinable

Consultations for Final Exam

- Please come to the consultations prepared
 - Do not ask questions like “Can you please explain Dynamic Programming from scratch?”.
- Don't try getting hints about the questions on final exam!
 - E.g., Is Kruskal's algorithm going to be on the exam?
- Don't ask how hard the exam is
 - It is, **on average**, easier than questions in the tutorial
 - It is designed to test your knowledge of the unit
 - It is designed to allow everyone a chance to get some marks, but not allow everyone to get full marks (i.e. a spread of difficulty)

Suggestions for preparation

- **Understand** how each algorithm works
- Practice writing pseudocode for each algorithm
- **Understand** its complexity analysis
- Don't confuse algorithms: Bellman-Ford vs Floyd-Warshall vs Ford-Fulkerson
 - Despite warning, every semester, students mix up algorithms losing all marks for the question
- Go over the early material and the week 11/12 material
- Go over the material which was not covered by assignments

Overview

- Topological Sort
 - Kahn's Algorithm
 - Depth First Search
- Design Principles (FIT2004 Summary)
- Final Exam etc.
- Review of all lecture material

Lecture 1

- correctness proof
- complexity recap
- recurrence relations
- proof by induction

Lecture 2

- intro to space complexity
- comparison costs
- stability
- selection sort analysis
- insertion sort analysis
- proof of lower bound for comparison based sorts
- count sort
- stable count sort
- radix sort
- recursive complexity (space and time)
- output sensitive time complexity

Lecture 3

- quicksort review
- partition out of place/in place/stable
- complexity analysis of quicksort (best/worst/average)
- kth order stats
- quickselect
- quickselect complexity
- median of medians (not examinable)

Lecture 4

- intro to DP
- Fibonacci
- coin change
- unbounded knapsack
- 0/1 knapsack
- edit distance
- constructing optimal solutions (finding coins)
- backtracking vs decision array

Lecture 5

- Hash tables
- direct addressing
- hashing/collision
- Birthday paradox
- collisions always occur
- ideal properties of a hash
- open hashing (chaining)
- closed hashing
- linear probing with deletion (lazy)
- primary clustering
- quadratic probing
- secondary clustering
- double hashing
- cuckoo hashing
- BST
- search
- insert
- delete
- worst case shape
- avl tree
- balance factor
- rebalancing
- complexity analysis of AVL

Lecture 6

- trie
- construction
- edge-node labels
- search
- nodes being arrays
- pros and cons
- properties
- suffix trie
- substring search
- lookup
- counting occurrences of substring using tree
- longest repeated substring using tree
- suffix tree
- suffix array
- querying SA
- $O(n)$ space SA
- longest repeated substring
- Construction of SA
- prefix doubling

Lecture 7

- BWT
- Last-First property
- justification of symbol clustering
- k-mers BWT inversion
- LF-mapping
- efficient inversion
- practice
- substring search + complexity

Lecture 8

- graph recap
- graph definition
- representation (adj matrix, adj list)
- BFS
- DFS
- some applications of BFS and DFS
- BFS for distances in unit weight
- dijkstras algorithm
- updating the heap vs double inserting
- proof of correctness
- stopppping early with single target
- recovering path

Lecture 9

- dijkstras with negative weights does not work
- negative cycles make shortest path meaningless
- bellman ford
- correctness
- unreachable cycles
- all pairs shortest paths
- floyd warshall
- correctness
- transitive closure

Lecture 10

- spanning tree
- minimum spanning tree
- general strategy of adding safe edges
- prims algorithm
- correctness
- kruskals algorithm
- correctness
- union find


Lecture 11

- Flow problem
- Flow network properties
- For Fulkerson algorithm
- Augmenting paths
- Complexity analysis
- Cuts
- Max-Flow / Min-Cut
- Proof of correctness

Lecture 12

- Kahn's algorithm
- Complexity
- DFS for topological sort
- Summary of unit
- Overview of exam
- Review of all lecture material

Coming Up Next



SWOT VAC